

LLM API Description — AI Studio Platform

LG Magazine (mybook.lgacademy.com) · request contract & per-task catalog for prompt registration

Prepared 2026-06-26 · Workspace ID 41539B272A... · Auth key No.1 (VGm...)

1. How the service calls AI Studio

The app authenticates with the workspace ID + API key to obtain a JWT, then issues one HTTPS POST per LLM task:

```
# Auth (cached JWT, auto-refresh)
POST {AISTUDIO_BASE_URL}/auth      # {ID: workspaceId, Password: apiKey} -> JWT

# Each task call
POST {AISTUDIO_BASE_URL}/genai/{API_CODE}/prompt/{PROMPT_INDEX}
Authorization: Bearer {JWT}
Body: {
  "empNo": "...",
  "historyId": null,
  "parameters": [
    { "key": "SYSTEM", "value": "<persona / instructions>" },
    { "key": "USER",   "value": "<task content>" }
  ]
}
# Response: generated text (+ token usage: inputTokens / outputTokens)
```

The persona and all task content are supplied by the app at call time in the SYSTEM and USER parameters. The registered prompt only needs to pass those two parameters through to the model — so a single registered prompt can serve many tasks.

2. What we need from the platform: two pass-through prompts

To run the dual-model design (heavy generation on a strong model, frequent one-liners on a light model), please register **two prompts**, each just forwarding {SYSTEM} + {USER} :

Prompt	Bind to model	Used for	Output size	Call share
PROMPT A	Claude Sonnet 4.6	Chapter 2/3 summaries & result-page (결과지) generation — quality-critical	large (~600–2,200 out tokens)	~20–25% of calls
PROMPT B	Gemini 2.5 Flash	L-OWL line-by-line reflections during the flow — light & frequent	small (~100–1,000 out tokens)	~75–80% of calls

Why split: the light one-liners are the bulk of the volume. Routing them to Gemini frees the strong model for the few heavy summaries, **lowering latency & cost and easing concurrency pressure.**

3. API task catalog

Each row is one task the app dispatches. A = Sonnet 4.6 prompt, B = Gemini 2.5 Flash prompt.

Task	Model	Purpose	Output	~Out tok	Freq / session
writeChapterArticle	A	Chapter 2/3 article / summary body	prose	~1,500	2–3
synthesizeStrength	A	Strength synthesis (result page)	structured prose	~2,200	1
synthesizeGrowthVision	A	Growth-vision synthesis (result page)	structured prose	~2,200	1
generateVisionDirections	A	Vision direction cards (result page)	JSON / cards	~800	1
generateJobTrendCards	A	Job-trend cards (result page)	JSON / cards	~600	1
generateTimeHorizon	A	Time-horizon section (result page)	structured	~500	1
writeEditorNote	A	Editor's note (intro/outro)	prose	~400	1–2
writeCoverHeadline	A	Cover headline	short line	~100	1
judgeBranch	B	Classify the answer to route the conversation	1 char (A/B/C/D)	~100	~10–15
reflectShort	B	L-OWL short reflection on an answer	1–2 lines	~1,000	~6–10
rephraseLight	B	Light rephrase of a line	1 line	~300	2–4
comfortReassure	B	Reassurance line	1–2 lines	~220	2–4
reflectPoetic	B	Poetic reflection across two stories	1–2 lines	~300	1–2
reflectValues	B	Reflection on chosen values	1–2 lines	~300	1
reflectStrength	B	Reflection on a strength	1–2 lines	~300	1

<code>extractKeyword</code>	B	Extract a keyword from an answer	1 word	~80	3–6
<code>observePattern</code>	B	Note a pattern across answers	1 line	~300	1–2

4. Capacity & sizing (per completed session)

	Sonnet 4.6 (Prompt A)	Gemini 2.5 Flash (Prompt B)
Calls per session	~8–10	~30–45
Call latency	15–25 s (long)	2–5 s (short)
Per-minute limit (LG-stated)	6,000,000 tokens/min	4,024 requests/min
Concurrency limit	~20 (observed) — <i>please confirm</i>	please provide

- Requests to the platform team:
- 1. Register **Prompt B** bound to **Gemini 2.5 Flash** (pass-through `{SYSTEM}` + `{USER}`) and share its **API code + prompt index**.
 - 2. Confirm the **concurrent-call** (동시 호출) **limit** for **both** Sonnet 4.6 and Gemini 2.5 Flash on this workspace.
 - 3. Keep **Prompt A** on **Sonnet 4.6** for the heavy tasks (current behaviour).

At ~3.5 light calls/min per active user, Gemini's 4,024 req/min covers ~1,000+ concurrent users — per-minute quota is not the constraint; **concurrency is the number we need**. The app enforces its own per-replica gate so it never exceeds whatever ceiling you confirm (`gate × replicas < limit`).